

Progress Report

Bardia Mojra

January 29, 2021

Robotic Vision Lab

The University of Texas at Arlington

1 To Do

- Implement basic estimation and generate a baseline for incremental advancement: Currently, I can train a basic model using the randomly generated data. I am currently working on resolving a visualization issue.
- Implement DOPE with added dropout before each layer to estimate variational Bayesian inference.
- Read Data Association and Localization of Classified Objects in Visual SLAM: in progress.
- Work on Pose Estimation Uncertainty paper.
- Implement PoseCNN, DOPE, and BayesOD.
- Look into domain randomization and adaptation techniques.
- Search for recent pose estimation survey papers: Found this paper, [1], on pose estimation.
- As I read more papers, go back and update Pose Estimation Survey paper.
- Finish Tensorflow tutorials, [2]: 4/9 chapters done.
- Learn to implement transfer learning.
- (On pause) Finish Docker tutorials, [3]: 4/18 chapters done.

2 Literature Reviews

2.1 Dropout as a Bayesian Approximation: Representing Model Uncertainty Deep Learning

In this paper [4], the authors introduce a mathematical framework that casts training deep neural nets with dropout layers as approximate Bayesian inference in deep Gaussian processes. They show that dropout layers in deep networks can minimize Kullback-Leibler divergence between an approximate distribution and the posterior of a deep Gaussian process, [5]. This is quite effective since dropout layers preserve Gaussian processes' variational distributions.

To make posterior weight distribution $p(\omega|D)$ intractable, columns are set to zero at random and we end up with $q(\omega)$, an approximate posterior weight distribution. Moreover, these layers are highly multi-modal and induces strong joint correlations over the rows of the network weight matrices W_i .

The objective KL divergence minimization is computed by approximating each term by Monte Carlo integration with a single sample, $\hat{\omega}_n \sim q(\omega)$. This allows for unbiased estimation of each posterior modality with $-\log p(y_n|x_n, \hat{\omega}_n)$ combined with the second summation term they obtained the following:

$$\mathcal{L}_{GP-MC} \propto \frac{1}{N} \sum_{n=1}^N \frac{-\log p(y_n|x_n, \hat{\omega}_n)}{\tau} + \sum_{l=1}^L \left(\frac{p_l l^2}{2\tau N} \|\mathbf{M}_l\|_2^2 + \frac{l^2}{2\tau N} \|\mathbf{m}_l\|_2^2 \right)$$

Where

$$E(y_n, \hat{y}(x_n, \hat{\omega}_n)) = -\log p(y_n|x_n, \hat{\omega}_n)/\tau$$

and approximate predictive distribution function represented by:

$$q(y^*|x^*) = \int p(y^*|x^*, \omega) q(\omega) d\omega$$

Which means this Monte Carlo estimation, *MC Dropout*, makes generalized predictions by averaging the weights of the network. This would be similar to T stochastic forward pass through the network and averaging the outputs but rather computationally more efficient.

For regression tasks, predictive log-likelihood estimate is given by:

$$\log p(y^*|x^*, X, Y) \approx \log \text{sumexp} \left(-\frac{1}{2} \tau \|y - \hat{y}_t\|^2 \right) - \log T - \frac{1}{2} \log 2\pi - \frac{1}{2} \log \tau^{-1}$$

In their experiments, they used a LeNet-5 with dropout layers added. For regression, MC Dropout with .1 or .2 probabilities (resulted in identical uncertainties) were applied before every weight layer and with TanH activation. For classification, they applied a dropout layer before the last fully connected inner-product layer. They use probability of .5 for dropout layers in classification tasks.

3 Progress

Following items are listed in order of priority:

- Object Pose Estimation with Uncertainty: I started a Tensorflow [6] project based on a simple implementation [7]. The code is written using Tensorflow Graphics module which is currently on supported by Google Colab, I am in the process of porting the code so it will use TensorBoard 3D instead. Next, I will expand the 3-layer network to five layers, train, test and compare the results. Later, I will add dropout layers for uncertainty analysis in regression.

I am currently restructuring the code into modules for model build, train, and testing. Other modules such as feature extraction can also be added to further enhance the performance.

- Implement PoseCNN, DOPE, and BayesOD.
- Look into domain randomization and adaptation techniques.
- Search for recent pose estimation survey papers: Found this paper, [1], on pose estimation.

4 Plans

Following items are listed in order of priority:

- Implement multiple object pose estimation with uncertainty estimation.
- Keep working on Bayesian Pose Estimation paper.
- ARIAC: For now, I will focus on implementing pose estimation and BayesOD implementations.
- Continue on UE4 tutorials.
- Pose Estimation Survey Paper Feedback: On hold, I am working on Bayesian Pose Estimation.
- Project Alpe with Nolan: On pause for right now.
- UR5e: Finish ROS Industrial tutorials.

5 2021 Goals and Target Journals/Conferences

- Submit a paper on pose estimation with uncertainty to ICIRS.
- Get comfortable with TensorFlow and related Python modules.
- Keep writing.

References

- [1] G. Du, K. Wang, S. Lian, and K. Zhao, “Vision-based robotic grasping from object localization, object pose estimation to grasp estimation for parallel grippers: a review,” *Artificial Intelligence Review*, pp. 1–58, 2020.
- [2] B. Planche and E. Andres, “Hands-on computer vision with tensorflow 2,” 2019.
- [3] G. Schenker, *Learn Docker – Fundamentals of Docker 19.x: Build, test, ship, and run containers with Docker and Kubernetes, 2nd Edition*. Packt Publishing, 2020.
- [4] Y. Gal and Z. Ghahramani, “Dropout as a bayesian approximation: Representing model uncertainty in deep learning,” 2016.
- [5] A. C. Damianou and N. D. Lawrence, “Deep gaussian processes,” 2013.
- [6] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G. S. Corrado, A. Davis, J. Dean, M. Devin, S. Ghemawat, I. Goodfellow, A. Harp, G. Irving, M. Isard, Y. Jia, R. Jozefowicz, L. Kaiser, M. Kudlur, J. Levenberg, D. Mané, R. Monga, S. Moore, D. Murray, C. Olah, M. Schuster, J. Shlens, B. Steiner, I. Sutskever, K. Talwar, P. Tucker, V. Vanhoucke, V. Vasudevan, F. Viégas, O. Vinyals, P. Warden, M. Wattemberg, M. Wicke, Y. Yu, and X. Zheng, “TensorFlow: Large-scale machine learning on heterogeneous systems,” 2015. Software available from tensorflow.org.
- [7] “graphics/6dof_alignment.ipynb at master · tensorflow/graphics.” https://github.com/tensorflow/graphics/blob/master/tensorflow_graphics/notebooks/6dof_alignment.ipynb. (Accessed on 01/28/2021).